

Predicting_Property_Maintenance_Fines

March 16, 2023

0.1 Project: Understanding and Predicting Property Maintenance Fines

0.2 Table of Contents

Introduction

Data Wrangling

Model selection training and evaluation

Conclusions

Introduction

[Blight violations](#) are issued by the city to individuals who allow their properties to remain in a deteriorated condition. Every year, the city of Detroit issues millions of dollars in fines to residents and every year, many of these fines remain unpaid. Enforcing unpaid blight fines is a costly and tedious process, so the city wants to know: how can we increase blight ticket compliance?

This project is focused on understanding when and why a resident might fail to comply with a blight ticket by training a model to predict blight ticket compliance in Detroit using `readonly/train.csv`. Using this model, return the probability that each corresponding ticket from `readonly/test.csv` will be paid.

0.2.1 Dataset Description

This project is based on a data challenge from the Michigan Data Science Team ([MDST](#)) who partnered with the City of Detroit to help solve one of the most pressing blight problem facing Detroit.

Each row in these two files corresponds to a single blight ticket, and includes information about when, why, and to whom each ticket was issued. The target variable is compliance, which is True if the ticket was paid early, on time, or within one month of the hearing date, False if the ticket was paid after the hearing date or not at all, and Null if the violator was found not responsible. Compliance, as well as a handful of other variables that will not be available at test-time, are only included in train.csv.

Data fields train.csv & test.csv

`ticket_id` - unique identifier for tickets

`agency_name` - Agency that issued the ticket

`inspector_name` - Name of inspector that issued the ticket

`violation_name` - Name of the person/organization that the ticket was issued to

`violation_street_number`, `violation_street_name`, `violation_zip_code` - Address where the violation

mailing_address_str_number, mailing_address_str_name, city, state, zip_code, non_us_str_code,
 ticket_issued_date - Date and time the ticket was issued
 hearing_date - Date and time the violator's hearing was scheduled
 violation_code, violation_description - Type of violation
 disposition - Judgment and judgement type
 fine_amount - Violation fine amount, excluding fees
 admin_fee - \$20 fee assigned to responsible judgments

state_fee - \$10 fee assigned to responsible judgments late_fee - 10% fee assigned to responsible
 judgments discount_amount - discount applied, if any clean_up_cost - DPW clean-up or graffiti
 removal cost judgment_amount - Sum of all fines and fees graffiti_status - Flag for graffiti violations

train.csv only

payment_amount - Amount paid, if any
 payment_date - Date payment was made, if it was received
 payment_status - Current payment status as of Feb 1 2017
 balance_due - Fines and fees still owed
 collection_status - Flag for payments in collections
 compliance [target variable for prediction]
 Null = Not responsible
 0 = Responsible, non-compliant
 1 = Responsible, compliant
 compliance_detail - More information on why each ticket was marked compliant or non-compliant

0.3 Evaluation

The predictions give the probability that the corresponding blight ticket will be paid on time. The evaluation metric for this project is the Area Under the ROC Curve (AUC). Model gives AUC score of above 0.75

Example:

```
ticket_id
284932    0.531842
285362    0.401958
285361    0.105928
285338    0.018572
...
376499    0.208567
376500    0.818759
369851    0.018528
Name: compliance, dtype: float32
```

Data Wrangling

```
[133]: import pandas as pd, numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Load the data files
train = pd.read_csv('train.
↳ csv',encoding='ISO-8859-1',low_memory=False,parse_dates=['ticket_issued_date',
↳ 'hearing_date', 'payment_date'])
test = pd.read_csv('test.csv',encoding='ISO-8859-1',low_memory=False,
↳ parse_dates=['ticket_issued_date', 'hearing_date'])
address = pd.read_csv('addresses.csv',encoding='ISO-8859-1',low_memory=False)
coord = pd.read_csv('latlons.csv',encoding='ISO-8859-1',low_memory=False)
```

Visually examine tables

```
[110]: train.sample(3)
```

```
[110]:      ticket_id      agency_name \
126148      153633  Buildings, Safety Engineering & Env Department
235889      270124  Buildings, Safety Engineering & Env Department
90156       115317  Buildings, Safety Engineering & Env Department

      inspector_name      violator_name  violation_street_number \
126148  Samaan, Neil J      HUTCHINS, OTTO      19304.0
235889  Samaan, Neil J      BISZCZANIK, MAREK      7425.0
90156   Samaan, Neil J  MISSIONARY, EMMANUEL      1271.0

      violation_street_name  violation_zip_code  mailing_address_str_number \
126148      HOOVER      NaN      1414.0
235889      DAVISON      NaN      6876.0
90156      OAKMAN BLVD      NaN      21143.0

      mailing_address_str_name      city  ... clean_up_cost judgment_amount \
126148      PO BOX      WARREN  ...      0.0      305.0
235889      MEADOWLAKE RD  BMFD TWP  ...      0.0      305.0
90156      PO BOX      DETROIT  ...      0.0      305.0

      payment_amount  balance_due  payment_date      payment_status \
126148      0.0      305.0      NaT  NO PAYMENT APPLIED
235889      0.0      305.0      NaT  NO PAYMENT APPLIED
90156      0.0      305.0      NaT  NO PAYMENT APPLIED

      collection_status  grafitti_status      compliance_detail \
126148  IN COLLECTION      NaN  non-compliant by no payment
235889  IN COLLECTION      NaN  non-compliant by no payment
90156   IN COLLECTION      NaN  non-compliant by no payment

      compliance
126148      0.0
235889      0.0
90156      0.0
```

[3 rows x 34 columns]

```
[111]: train.shape
```

```
[111]: (250306, 34)
```

```
[112]: address.sample(3)
```

```
[112]:      ticket_id      address
152790    182446    1975 webb, Detroit MI
85121     110178    5418 iroquois, Detroit MI
163791    193718    2082 vinewood, Detroit MI
```

```
[113]: coord.sample(3)
```

```
[113]:      address      lat      lon
30389    4867 avery, Detroit MI  42.351081 -83.081252
115680    1155 lenore, Detroit MI  42.421118 -83.281344
106383    1524 military, Detroit MI  42.313517 -83.104955
```

```
[114]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250306 entries, 0 to 250305
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ticket_id                             250306 non-null  int64
1   agency_name                           250306 non-null  object
2   inspector_name                        250306 non-null  object
3   violator_name                         250272 non-null  object
4   violation_street_number               250306 non-null  float64
5   violation_street_name                 250306 non-null  object
6   violation_zip_code                    0 non-null      float64
7   mailing_address_str_number            246704 non-null  float64
8   mailing_address_str_name              250302 non-null  object
9   city                                 250306 non-null  object
10  state                                 250213 non-null  object
11  zip_code                             250305 non-null  object
12  non_us_str_code                       3 non-null      object
13  country                              250306 non-null  object
14  ticket_issued_date                   250306 non-null  datetime64[ns]
15  hearing_date                         237815 non-null  datetime64[ns]
16  violation_code                        250306 non-null  object
17  violation_description                 250306 non-null  object
18  disposition                           250306 non-null  object
```

19	fine_amount	250305	non-null	float64
20	admin_fee	250306	non-null	float64
21	state_fee	250306	non-null	float64
22	late_fee	250306	non-null	float64
23	discount_amount	250306	non-null	float64
24	clean_up_cost	250306	non-null	float64
25	judgment_amount	250306	non-null	float64
26	payment_amount	250306	non-null	float64
27	balance_due	250306	non-null	float64
28	payment_date	41113	non-null	datetime64[ns]
29	payment_status	250306	non-null	object
30	collection_status	36897	non-null	object
31	grafitti_status	1	non-null	object
32	compliance_detail	250306	non-null	object
33	compliance	159880	non-null	float64

dtypes: datetime64[ns](3), float64(13), int64(1), object(17)

memory usage: 64.9+ MB

```
[115]: train.isnull().sum()
```

```
[115]: ticket_id          0
agency_name             0
inspector_name          0
violator_name           34
violation_street_number  0
violation_street_name    0
violation_zip_code      250306
mailing_address_str_number 3602
mailing_address_str_name  4
city                    0
state                   93
zip_code                 1
non_us_str_code         250303
country                  0
ticket_issued_date      0
hearing_date            12491
violation_code           0
violation_description    0
disposition              0
fine_amount              1
admin_fee                0
state_fee                0
late_fee                 0
discount_amount          0
clean_up_cost            0
judgment_amount          0
payment_amount           0
```

```

balance_due          0
payment_date         209193
payment_status       0
collection_status    213409
grafitti_status      250305
compliance_detail    0
compliance           90426
dtype: int64

```

0.3.1 Strategize wrangling path by defining cleaning process on train table.

Quality

- Missing demographic information (payment_date, collection_status, grafitti_status, compliance, violator_name, violation_zip_code, mailing_address_str_number, mailing_address_str_name, state, zip_code, non_us_str_code, hearing_date contact columns) *(can't clean yet)*

Tidiness

- Columns with all entries being zero should be removed.
- Columns with the same values should be removed as they are uncorrelated with target variable.
- Columns with total unique values less than 10% of entries (<250) should be converted into categorical data to reduce memory usage to reduce memory usage.
- Remove columns with missing value % of more than 50%.
- Join the address table to train and test tables to expand features.
- With address joined, now remove features that can be replaced with address, such as :

```
['violator_name', 'violation_street_number', 'violation_street_name', 'mailing_address_str_number', 'mailing_address_str_name', 'state', 'zip_code', 'country', 'address', 'city']
```

- Reduce the features even further, by summing the amount payables into one.
- drop missing values of ['lat', 'lon', 'total_amt_pay'] from the train dataset
- Replace ticket issue data and the hearing date with the time gap between them.
- Now remove not too important featured and make string features from string categories

```
['inspector_name', 'violation_code', 'violation_description', 'payment_amount', 'balance_due', 'payment_status', 'compliance_detail']
```

- taking only non-NaN values for training
- trime the train data to have only the columns available in the test data

```
[116]: # now we remove columns and rows with all entries being EMPTY
train.dropna(how='all',axis=1, inplace=True)
train.dropna(how='all',axis=0, inplace=True)
```

```
[117]: # Remove columns with the same values they are independent/non-correlated to/
      ↪ from target values
independent = []
```

```

for i in range(len(train.columns)):
    if len(train[train.columns[i]].unique())==1:
        independent.append(train.columns[i])

train.drop(independent,axis=1,inplace=True)
test.drop(independent,axis=1,inplace=True)

```

[118]: *# we see that there are a lot of columns with total unique values less than 250.*
→Thus we can convert them into categorical data to reduce memory usage
to reduce memory usage we convert columns with < than 250 entries to
→categorical data

```

for i in range(len(train.columns)):
    if len(train[train.columns[i]].unique())<250:
        train[train.columns[i]] = train[train.columns[i]].astype('category')

```

[119]: *# now lets see the missing number ratio in the data set*
total_null = train.isnull().sum().sort_values(ascending=False)
per = train.isnull().count().sort_values(ascending=False)

[120]: *# now i remove columns with missing value percentage of more than 50%*

```

high_mssing_data = pd.concat([total_null,total_null/per],
    →keys=['Total_nulls','percentage_nulls'],axis=1)
high_missing_values = high_mssing_data[high_mssing_data['percentage_nulls']>0.
    →5].index
train.drop(high_missing_values,axis=1,inplace=True)

```

[121]: *# Now we join the address to train and test data*
address = address.merge(coord,how='inner',left_on='address',right_on='address')
train = train.
 →merge(address,how='left',left_on='ticket_id',right_on='ticket_id')\
 .set_index('ticket_id')
test = test.merge(address,how='left',left_on='ticket_id',right_on='ticket_id')\
 .set_index('ticket_id')

[122]: *# now we reduce the features that can be replaced by the lat and lon*
latlon_replaced = ['violation_name',
 'violation_street_number', 'violation_street_name',
 'mailing_address_str_number', 'mailing_address_str_name',
 'state', 'zip_code', 'country','address','city']
train.drop(latlon_replaced, axis=1,inplace=True)

[123]: *# Now we reduce the features even further, by suming the amount payables into*
→one

```

train['total_amt_pay'] =
    ↳train[['fine_amount','admin_fee','state_fee','late_fee']].sum(axis=1).
    ↳subtract(train['discount_amount'].astype(np.float64))
test['total_amt_pay'] =
    ↳test[['fine_amount','admin_fee','state_fee','late_fee']].sum(axis=1).
    ↳subtract(test['discount_amount'].astype(np.float64))
drop_payments =
    ↳['fine_amount','admin_fee','state_fee','late_fee','discount_amount']
train.drop(drop_payments,axis=1, inplace=True)

```

```

[124]: # drop missing values of ['lat','lon','total_amt_pay'] from the train dataset
    ↳but since its not allowed in the test set,we replace it with the mean
train.dropna(subset = ['lat','lon','total_amt_pay'],inplace=True)
test['lat'].fillna(test.lat.mean(),inplace=True)
test['lon'].fillna(test.lon.mean(),inplace=True)

```

```

[125]: # Now we find the time gap between the ticket issue data and the hearing date
train['time_delta'] = (train['hearing_date'] - train['ticket_issued_date']).dt.
    ↳days
test['time_delta'] = (test['hearing_date'] - test['ticket_issued_date']).dt.
    ↳days
drop_timedelta = ['hearing_date','ticket_issued_date']
train.drop(drop_timedelta,axis=1, inplace=True)
test.drop(drop_timedelta,axis=1, inplace=True)

```

```

[126]: # Replace the missing values in the time delta column with the mode
train['time_delta'].fillna(73, inplace=True)
test['time_delta'].fillna(73,inplace=True)

```

```

[127]: # Now remove not too important featured and make string features from string
    ↳categories 'disposition','agency_name'
further_drop = ['inspector_name', 'violation_code','violation_description',
    ↳'payment_amount', 'balance_due','payment_status',
    ↳'compliance_detail']

train.drop(further_drop,axis=1, inplace=True)
string_features = ['disposition','agency_name']
train = pd.get_dummies(train,columns = string_features,drop_first=True)
test = pd.get_dummies(test,columns = string_features,drop_first=True)

```

```

[128]: # taking only non-NaN values for training
train = train[(train['compliance']==0) | (train['compliance']==1) ]

```

```

[129]: # trime the train data to have only the columns available in the test data
y = train['compliance']
X = train.drop('compliance',axis=1)

```



```

train_feature_set = set(X)
for feature in set(X):
    if feature not in test:
        train_feature_set.remove(feature)

train_features = list(train_feature_set)
X_train = X[train_features]
test      = test[train_features]

# X_train, y, test,

```

Model selection, training and evaluation

```

[53]: from sklearn.model_selection import train_test_split
      ↪      ,GridSearchCV
      from sklearn.preprocessing import MinMaxScaler, StandardScaler
      from sklearn.dummy import DummyClassifier
      from sklearn.metrics import roc_auc_score
      from sklearn.neighbors import KNeighborsRegressor
      from sklearn.linear_model import RidgeClassifier
      from sklearn.svm import SVC
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.naive_bayes import GaussianNB
      from sklearn.ensemble import RandomForestClassifier\
      ,RandomForestRegressor\
      ,GradientBoostingClassifier
      import time

```

```

[54]: # divide train test split for model selection
      x_train,x_test,y_train,y_test= train_test_split(X_train,y,test_size=0.
      ↪4,random_state=40)
      X_train.shape,y.shape

```

```

[54]: ((159878, 11), (159878,))

```

```

[61]: # implement dummy classifier as base model, using most frequent data as
      ↪prediction
      start = time.time()
      dummy_clf = DummyClassifier(strategy='most_frequent').fit(x_train,y_train)
      y_pred = dummy_clf.predict(x_test)
      print('Runtime: {} \nROC_score : {}'.format(time.time()- start,\
      roc_auc_score(y_test,y_pred)))

```

```

Runtime: 0.01413583755493164
ROC_score : 0.5

```

```
[65]: # implementing KNN classifier
start = time.time()
KN_clf = KNeighborsRegressor()
param_values = {'n_neighbors':[1,3,5,7,9]}
grid_clf = GridSearchCV(KN_clf, param_grid=param_values,scoring='roc_auc').
    ↪fit(x_train,y_train)
print('Best parameter is: {} \n Best ROC_score is: {}'.format(grid_clf.
    ↪best_params_,grid_clf.best_score_))
print('Runtime: {}'.format(time.time()-start))
```

Best parameter is: {'n_neighbors': 9}
 Best ROC_score is: 0.7494481978035308
 Runtime: 29.89334988594055

```
[68]: # implenting Ridge classifier
start = time.time()
R_clf = RidgeClassifier()
R_clf.get_params().keys()
param_values = {'alpha':[0, 1, 10, 20, 50, 100, 1000]}
grid_clf = GridSearchCV(R_clf, param_grid=param_values,scoring='roc_auc').
    ↪fit(x_train,y_train)
print('Best parameter is {} \n Best ROC_score is {}'.format(grid_clf.
    ↪best_params_,grid_clf.best_score_))
print('Runtime: {}'.format(time.time()-start))
```

/Users/air/opt/anaconda3/lib/python3.9/site-
 packages/sklearn/linear_model/_ridge.py:157: LinAlgWarning: Ill-conditioned
 matrix (rcond=1.03558e-22): result may not be accurate.
 return linalg.solve(A, Xy, sym_pos=True, overwrite_a=True).T
 /Users/air/opt/anaconda3/lib/python3.9/site-
 packages/sklearn/linear_model/_ridge.py:157: LinAlgWarning: Ill-conditioned
 matrix (rcond=9.24063e-23): result may not be accurate.
 return linalg.solve(A, Xy, sym_pos=True, overwrite_a=True).T
 Best parameter is {'alpha': 0}
 Best ROC_score is 0.7723439974165995
 Runtime: 3.734174966812134

/Users/air/opt/anaconda3/lib/python3.9/site-
 packages/sklearn/linear_model/_ridge.py:157: LinAlgWarning: Ill-conditioned
 matrix (rcond=5.47116e-23): result may not be accurate.
 return linalg.solve(A, Xy, sym_pos=True, overwrite_a=True).T

```
[72]: # implenting random forest regressor
start = time.time()
RF_clf = RandomForestRegressor()
param_values = {'max_depth':[1,3,5,7,9,11,13,15,17,18,21]}
# X_train = MinMaxScaler().fit_transform(X_train)
```

```
# # testt = MinMaxScaler().fit_transform(test)
grid_clf = GridSearchCV(RF_clf, param_grid=param_values,scoring='roc_auc').
    ↪fit(x_train,y_train)
print('Best parameter is {}\\nBest ROC_score is {}'.format(grid_clf.
    ↪best_params_,grid_clf.best_score_))
print('Runtime: {}'.format(time.time()-start))
```

Best parameter is {'max_depth': 15}
 Best ROC_score is 0.8197865894627888
 Runtime: 1321.5681660175323

```
[73]: # implementing support vector machine classifier
start = time.time()
SV = SVC(kernel = 'rbf',C = 0.01)
param_value = {'gamma': [0.01,0.1,1,10,100,]}
grid_clf = GridSearchCV(SV,param_grid= param_value ,scoring = 'roc_auc').
    ↪fit(x_train,y_train)
print('Best parameter is {}\\nBest ROC_score is {}'.format(grid_clf.
    ↪best_params_,grid_clf.best_score_))
print('Runtime: {}'.format(time.time()-start))
```

Best parameter is {'gamma': 100}
 Best ROC_score is 0.6964877516589298
 Runtime: 4802.026435136795

Random forest regressor is selected

```
[130]: def blight_model(X_train, y, test):
    #import necessary models to train the data
    from sklearn.ensemble import RandomForestRegressor

    RF_clf = RandomForestRegressor(max_depth=6).fit(X_train,y)
    y_pred = RF_clf.predict(test)
    test['compliance'] = y_pred
    return test.compliance
```

```
[131]: blight_model(X_train, y, test)
```

```
[131]: ticket_id
284932    0.063016
285362    0.013487
285361    0.068180
285338    0.083102
285346    0.087070
...
376496    0.013312
376497    0.013312
376499    0.070316
```

```
376500    0.070316
369851    0.399647
Name: compliance, Length: 61001, dtype: float64
```

```
[ ]:
```